



SCALE THE PROVEN

Client Configuration Template

Every workflow parameter as config — accounts, consent basis, safety keywords, pricing source, voice, and launch state. NEEDS-INPUT blocks launch.

Client Configuration Template

One file per client, copied from `clients/_template/`. Every workflow parameter lives here — nothing client-specific is ever hardcoded in prompts, SOPs, or n8n workflows. A `NEEDS-INPUT` left in any critical field means the client is **not launchable**.

Used by the [Client Onboarding SOP](#) at Gate 2 (configuration) and Gate 5 (launch state). Real client files are kept private — never committed to a shared repo.

```
# Client AI Brain configuration – ONE file per client, copied from _template.
# Every workflow parameter lives here; nothing client-specific is hardcoded
# in prompts, SOPs, or n8n workflows. NEEDS-INPUT in a critical field = not launchable.
# Real client dirs are gitignored (they contain business-sensitive data);
# only this template is committed. See sops/client-onboarding.md.

client:
  code: NEEDS-INPUT          # short slug, e.g. sheren-ph (tags audit_log, Langfuse,
  folders)
  business_name: NEEDS-INPUT # exactly as it should appear in messages
  trade: NEEDS-INPUT         # hvac | plumbing | gutter | electrical | multi
  timezone: America/Detroit
  onboarded: NEEDS-INPUT     # YYYY-MM-DD
  package: NEEDS-INPUT       # leakage-audit | pilot-45d | build-90d | retainer

accounts:                    # THEIR accounts – isolation rule, see Gate 1
  ghl_location_id: NEEDS-INPUT
  n8n_instance: NEEDS-INPUT  # url/host of THEIR n8n workspace
  langfuse_project: NEEDS-INPUT
  llm_keys: NEEDS-INPUT      # theirs | censeoai-tagged (per the deal)

contact_rules:
  business_hours: "Mon-Fri 08:00-17:00, Sat 09:00-13:00" # adjust
  send_window: "08:00-21:00" # recipient-local; hard-coded gate in n8n
  max_agent_messages_without_reply: 3
  tcpa_basis: NEEDS-INPUT    # e.g. "inbound call from consumer's own number; single
  contextual text-back; STOP honored"
  dnc_list_imported: false   # must be true before launch
  opt_out_text: "Reply STOP to opt out"

escalation:
```

```

phone: NEEDS-INPUT          # on-call human, confirmed in writing by owner
approver: NEEDS-INPUT       # HITL approver: name + channel (their office mgr |
censeoai)
safety_keywords:           # adjust per trade (Gate 2)
  - gas
  - leak
  - smell
  - smoke
  - carbon monoxide
  - flood
  - sparking
  - no heat
  - no AC
always_human:              # from the owner walk-through (Gate 3)
  - NEEDS-INPUT

pricing:
  source: NEEDS-INPUT       # ghl-custom-fields | postgres-pricing-table
  populated: false          # true only when their real price book is loaded
  # agents NEVER quote prices not found in the source above

voice:
  review_link: NEEDS-INPUT
  samples: |                # 3-5 real messages from their front desk (Gate 3)
    NEEDS-INPUT

followup:                  # workflow 03 offsets
  satisfaction_check: same-evening
  review_request_days: 2
  checkin_days: 30

launch:
  hitl_gate: true           # stays true until flip checklist passes (Gate 5)
  flip_checklist:
    module_a_recognition_current: false
    tool_validation_id: null
    approved_sends: 0       # needs 30+ at >=95% unmodified
  test_run:                # Gate 4 – fill after the 7-case test passes
    date: null
    run_by: null
    results: null

```

Field notes

- **contact_rules.tcpa_basis** — write the actual consent story in plain English; "we have consent" is not a basis.
- **pricing.populated** — stays `false` until the client's real price book is loaded. Agents never quote a price the source can't return; an empty table means "our team will confirm," which is correct behavior.
- **escalation.always_human** — comes from the one-hour owner session (Gate 3), in the owner's words.
- **launch.flip_checklist** — all three conditions must be true before the HITL gate comes off, per client, every time: current Module A recognition, a recorded tool-validation ID, and 30+ approved sends at ≥95% unmodified.
- **launch.test_run** — filled only after the 7-case test passes with THIS client's config (Gate 4).

CenseoAI · Scale the Proven · www.CenseoAI.ai | Source: [censeo-ai-brain/clients/_template/config-template-doc.md](#) · generated 2026-07-25 · [links open on GitHub](#)